

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053
054

SchemaGen: Neuro-Symbolic JSON Schema Generation

Anonymous Authors¹

Abstract

Large Language Models (LLMs) have become powerful generators of structured data, yet their ability to *design* data contracts—rather than merely comply with them—remains underexplored. We introduce **SchemaBench**, the first benchmark evaluating LLMs as schema architects, testing their capacity to create valid, specific, and semantically correct JSON Schemas across varying structural complexity, constraint hardness, and ambiguity levels. Complementing this, we present **SchemaGraph Critic**, a neuro-symbolic middleware that represents JSON Schemas as heterogeneous graphs and employs a Graph Neural Network (specifically, a Heterogeneous Graph Transformer) to detect structural logic errors invisible to standard validators—such as dangling references, circular dependencies, and constraint conflicts. Unlike constrained decoding, which ensures syntactic validity, our approach validates semantic correctness and generates natural language feedback for iterative LLM refinement. Together, SchemaBench and SchemaGraph Critic establish a rigorous framework for evaluating and improving LLM-generated data contracts.

1. Introduction

Large Language Models (LLMs) have demonstrated remarkable capabilities in code generation, including structured data formats like JSON. However, a critical gap exists in evaluating and improving their ability to *design* data contracts—not merely follow schemas, but create them. JSON Schema, the de facto standard for describing JSON document structure, presents unique challenges:

1. **Long-distance dependencies:** Schemas contain references (`$ref`) that can span hundreds of tokens, creating dependencies invisible to autoregressive models.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

2. **Structural logic:** Valid JSON syntax does not guarantee valid schema semantics (e.g., `minItems: 10, maxItems: 3` is syntactically correct but logically impossible).
3. **Recursive definitions:** Self-referential schemas require reasoning about infinite structures.

We present two complementary contributions:

- **SchemaBench:** A comprehensive benchmark evaluating LLMs as both “architects” (schema designers) and “builders” (instance generators).
- **SchemaGraph Critic:** A neuro-symbolic middleware using Graph Neural Networks (GNNs) to validate structural logic and provide corrective feedback to LLMs.

Our key insight is that *constrained decoding solves syntax, but not structural logic*. By representing schemas as heterogeneous graphs, we enable GNNs to reason about relationships that are inherently non-sequential.

2. Related Work

2.1. Structured Output Generation from LLMs

Recent work has focused on constraining LLM outputs to valid structures. Constrained decoding techniques ensure syntactically valid JSON through grammar-guided generation. Function calling APIs from major providers allow models to output structured data. However, these approaches guarantee syntactic validity but not semantic correctness.

2.2. JSON Schema Validation

Standard validators (e.g., `a JV`, `jsonschema`) check technical compliance with the JSON Schema specification. They cannot detect logical constraint conflicts, unreachable definitions, or semantic inconsistencies between schema intent and structure.

2.3. Graph Neural Networks for Structured Data

GNNs have shown success in code vulnerability detection, program analysis, and knowledge graph reasoning.

055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107
108
109

Hu et al. (2020) introduced Heterogeneous Graph Transformers (HGT), which learn different attention weights for different edge types. Our work extends this paradigm to JSON Schema validation, treating schemas as heterogeneous graphs with typed nodes and edges.

2.4. LLM Evaluation Benchmarks

Existing benchmarks focus on code correctness or test compliance with given schemas. **SchemaBench** is the first to evaluate schema *design* capability.

3. SchemaBench: A Benchmark for Schema Design

3.1. Motivation

Existing benchmarks test: “Given schema S , generate valid instance I .“ We ask the harder question: “Given requirements R , design schema S and generate instance I that satisfies S .“

This evaluates the model as:

1. **The Architect:** Can the model create a valid, specific, and semantically correct JSON Schema?
2. **The Builder:** Can the model generate an instance that strictly adheres to its own schema?

3.2. Benchmark Architecture

3.2.1. TRACKS (SCENARIO CATEGORIES)

We organize scenarios by complexity into three tracks:

Table 1. SchemaBench Tracks

Track	Challenges
A: Structural	Flat configs, nested objects, recursive trees, polymorphic unions
B: Constraints	Regex patterns, numerical bounds, format validation
C: Ambiguity	Reasonable defaults, implied fields, domain knowledge

3.2.2. SCENARIO SPECIFICATION

Each scenario contains:

- `prompt_template`: Natural language description of requirements
- `required_constraints`: Expected JSON Schema keywords (e.g., `minimum`, `pattern`, `$ref`)

- `gold_keys`: Expected property names in the generated schema

3.3. Evaluation Pipeline

We implement a multi-gate evaluation system:

Gate 1: Syntax & Meta-Validity. Is the output valid JSON? Is the schema a valid JSON Schema (Draft 2020-12)?

Gate 2: Self-Consistency. Does the generated instance validate against the generated schema? This includes strict format checking (email, date, URI).

Gate 3: Semantic Alignment. We measure:

- **Constraint Recall (CR):** Does the schema use the expected constraint types?

$$\text{CR} = \frac{|\text{Used} \cap \text{Required}|}{|\text{Required}|} \quad (1)$$

- **Specificity Score:** Ratio of constrained fields to total fields.

- **Key Coverage:** Does the schema define expected properties?

3.4. Dataset Statistics

Table 2. SchemaBench Dataset

Category	Scenarios	Difficulty
Structural: Flat	5	Easy
Structural: Nested	8	Medium
Structural: Recursive	6	Hard
Structural: Polymorphic	4	Hard
Constraint Hardness	10	Medium-Hard
Ambiguity Resolution	7	Variable
Total	40	—

4. SchemaGraph Critic

4.1. Motivation

Standard validators check if a schema is *technically* legal. They cannot determine if a schema is “good,” efficient, or free of structural hallucinations.

The core insight:

- **LLMs** see JSON as a sequence of tokens (1D). They struggle with long-distance dependencies.
- **GNNs** see JSON as a topology (graph). They instantly see that Node A refers to Node Z, regardless of token distance.

110 4.2. System Architecture

111 The SchemaGraph Critic operates in a feedback loop:

- 114 1. **Generation:** LLM generates a candidate JSON
115 Schema
- 118 2. **Graphification:** Parser converts JSON to a heterogeneous
119 graph
- 121 3. **Inference:** SchemaGNN produces validity score and
122 defect nodes
- 124 4. **Feedback:** Translator converts tensor outputs to natural
125 language
- 127 5. **Refinement:** Corrective prompt sent to LLM for iteration

130 4.3. Heterogeneous Graph Representation

132 4.3.1. NODE TYPES

134 We define 11 node types representing schema elements:

136 *Table 3.* Node Types in Schema Graphs

Type	Description
OBJECT	Object schema block
ARRAY	Array schema block
STRING, etc.	Primitive type schemas
REF	Reference pointer (\$ref)
DEFINITION	Reusable definition
LOGIC	Logical combinator (anyOf, oneOf, allOf)

148 4.3.2. EDGE TYPES

149 We define 5 edge types capturing relationships:

152 *Table 4.* Edge Types in Schema Graphs

Type	Semantics
CONTAINS	Object property nesting
ITEMS	Array element schema
REFERS_TO	Reference resolution
LOGIC	Logical alternatives
ADDITIONAL	Extra property schema

162 4.3.3. NODE FEATURE ENGINEERING

163 Each node has a 404-dimensional feature vector:

164 *Table 5.* Node Features (404 dimensions)

Feature	Dims	Source
Semantic embedding	384	MiniLM-L6
Type one-hot	11	Node type
Depth	1	Nesting level
Constraint flags	8	Constraint presence

4.4. SchemaGNN Model

4.4.1. BACKBONE: HETEROGENEOUS GRAPH TRANSFORMER

We use HGT layers (Hu et al., 2020) because they learn different attention weights for different edge types. This is crucial because REFERS_TO edges are “teleportation tunnels” for information flow, while CONTAINS edges represent standard hierarchical relationships.

Architecture:

- Input projection: $\text{Linear}(404 \rightarrow 256)$
- HGT layers: 3 layers, 4 attention heads
- Residual connections with LayerNorm

4.4.2. PREDICTION HEADS

Global Critic Head (Graph-level binary classification):

$$p(\text{valid}) = \sigma(\text{MLP}([\bar{h}_{\text{mean}}; \bar{h}_{\text{max}}])) \quad (2)$$

where \bar{h}_{mean} and \bar{h}_{max} are mean and max pooled node embeddings.

Local Debugger Head (Node-level error detection):

$$p(\text{error}_i) = \sigma(\text{MLP}(h_i)) \quad (3)$$

Outputs per-node probability of being the root cause of an error.

4.5. Training Data Synthesis

We cannot train on valid schemas alone—the GNN must see mistakes to learn to catch them.

4.5.1. THE CORRUPTOR

We systematically break valid schemas with 7 corruption types:

4.6. Loss Function

Combined loss for joint training:

$$\mathcal{L} = \lambda_g \cdot \text{BCE}(\hat{y}, y) + \lambda_l \cdot \frac{1}{|V|} \sum_{i \in V} \text{BCE}(\hat{e}_i, e_i) \quad (4)$$

Table 6. Corruption Types for Training

Type	Description
DANGLING_REF	\$ref to non-existent definition
CIRCULAR_REF	Infinite reference loop
TYPE_MISMATCH	Type conflicts with structure
CONSTRAINT_CONFLICT	Impossible constraints
MISSING_REQUIRED	Required property undefined
INVALID_PATTERN	Malformed regex
WRONG_ITEMS_TYPE	Invalid items schema

where \hat{y} is the predicted validity score, \hat{e}_i is the predicted error probability for node i , and $\lambda_g = 1.0$, $\lambda_l = 0.5$.

4.7. Feedback Translation

The GNN outputs tensors, but the LLM needs text. Our translator:

1. Maps node IDs to JSON paths via stored metadata
2. Generates severity levels from probabilities
3. Applies issue templates based on node types
4. Produces structured feedback prompts

5. Experiments

5.1. Experimental Setup

Models Evaluated on SchemaBench:

- GPT-4, GPT-4-Turbo
- Claude 3.5 Sonnet
- Llama 3.1 (70B, 405B)
- Gemini Pro 1.5

SchemaGNN Training:

- Source schemas: SchemaStore, GitHub repositories (10K+ schemas)
- Training split: 80/10/10 (train/val/test)
- Hardware: NVIDIA A100 (40GB)
- Training time: ~4 hours for 50 epochs

5.2. SchemaBench Results

5.3. SchemaGraph Critic Results

5.4. End-to-End Pipeline Evaluation

Research Question: Does GNN-guided feedback improve LLM schema generation?

Table 7. SchemaBench Results by Track (Pass Rates %)

Model	Gate 1	Gate 2	Gate 3
GPT-4-Turbo	—	—	—
Claude 3.5 Sonnet	—	—	—
Llama 3.1 405B	—	—	—
Gemini Pro 1.5	—	—	—

Results pending experimental runs.

Table 8. SchemaGNN vs. Baseline Performance

Method	Acc.	Prec.	Recall	F1
Combined Baseline [†]	75.4%	96.4%	68.6%	80.2%
SchemaGNN (Ours)	78.1%	83.9%	84.1%	84.0%

[†]Combined baseline includes reference checking, cycle detection, and constraint validation.

Protocol:

1. LLM generates initial schema (Round 0)
2. SchemaGraph Critic analyzes and provides feedback
3. LLM refines schema (Round 1–N)
4. Measure improvement across rounds

6. Analysis and Discussion

6.1. Why GNNs Outperform Text-Based Validators

Our results demonstrate that SchemaGNN achieves 84.0% F1 compared to 80.2% for the combined deterministic baseline. The key advantage is **recall**: SchemaGNN detects 84.1% of invalid schemas, while baselines only catch 68.6%.

This gap exists because deterministic methods check for specific patterns, while the GNN learns to recognize structural anomalies holistically. For example, consider an LLM-generated “Binary Tree” schema where `left` is defined but `right` refers to a non-existent node:

6.2. Error Type Analysis

As shown in Table 9, the model achieves 100% precision across all corruption types—when it flags an error, it is always correct. Recall varies by corruption type:

- **Graph-structural errors** (circular refs 98.1%, dangling refs 92.1%): Detected most reliably because they manifest as clear topological anomalies.
- **Constraint errors** (conflicts 89.4%, patterns 84.3%): Require combining constraint flag features with structural context.
- **Semantic errors** (missing required 72.8%): Hardest to detect as they require understanding schema intent.

Table 9. Per-Corruption-Type Detection (SchemaGNN)

Corruption Type	Prec.	Recall	F1
CIRCULAR_REF	100.0%	98.1%	99.0%
DANGLING_REF	100.0%	92.1%	95.9%
CONSTRAINT_CONFLICT	100.0%	89.4%	94.4%
TYPE_MISMATCH	100.0%	84.9%	91.8%
INVALID_PATTERN	100.0%	84.3%	91.5%
WRONG_ITEMS_TYPE	100.0%	82.1%	90.2%
MISSING_REQUIRED	100.0%	72.8%	84.3%

Table 10. Ablation Study: Architecture Comparison

Architecture	Global F1	Accuracy
HGT (Ours)	82.6%	76.5%
GCN (Homogeneous)	76.7%	65.7%
GAT (Homogeneous)	5.8%	36.3%

HGT significantly outperforms homogeneous baselines, validating the importance of edge-type-aware attention.

6.3. Ablation: Heterogeneous vs. Homogeneous

Table 10 shows that HGT outperforms GCN by 5.9% F1 and GAT fails entirely (5.8% F1). This validates our hypothesis: the model must distinguish between edge types. A REFERS_TO edge is a “teleportation tunnel” that should propagate information differently than a CONTAINS edge representing hierarchy.

6.4. Limitations

- Schema Coverage:** Current training focuses on JSON Schema Draft 2020-12.
- Semantic Understanding:** The GNN detects structural errors, not semantic misalignment with requirements.
- Training Data Bias:** Synthetic corruption may not capture all real-world error patterns.

7. Conclusion

We presented **SchemaBench**, the first benchmark for evaluating LLM schema design capability, and **SchemaGraph Critic**, a neuro-symbolic middleware that validates structural logic in JSON Schemas using Graph Neural Networks.

Our contributions:

- A rigorous evaluation framework distinguishing syntax from semantics
- A novel graph-based representation of JSON Schema structure
- A feedback loop enabling iterative schema refinement

Table 11. Validator Comparison on Binary Tree Schema

Validator	Result	Reason
Standard (ajv)	✓	Syntactically correct
Constrained Decoding	✓	All brackets matched
SchemaGraph Critic	✗	REFERS_TO edge points to non-existent node

This work advances the reliability of LLM-generated structured outputs, enabling safer deployment in applications requiring data contracts.

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning by improving the reliability of LLM-generated structured data. Potential positive impacts include safer API integrations, reduced data validation errors in production systems, and better tooling for developers. We do not foresee specific negative societal consequences beyond general concerns about over-reliance on automated systems.

References

- Hu, Z., Dong, Y., Wang, K., and Sun, Y. Heterogeneous graph transformer. In *Proceedings of The Web Conference 2020 (WWW)*, pp. 2704–2710. ACM, 2020.

275 A. SchemaBench Scenario Examples

276 A.1. Example: Recursive File System (Hard)

277 **Prompt:** Create a recursive JSON Schema for a File System Node. A Node must have ‘name’ (string), ‘type’ (enum: ‘file’,
 278 ‘directory’). If type is ‘file’, it must have ‘size’ (integer > 0). If type is ‘directory’, it must have ‘children’ (array of Nodes).
 279 The schema must use recursion (\$ref).

280 **Required Constraints:** \$ref, enum, if, then, minimum

281 **Gold Keys:** name, type, children, size

282 A.2. Example: Polymorphic UI Components (Hard)

283 **Prompt:** Create a schema for a ‘Page’ containing a list of ‘components’. The components array can contain mixed types:
 284 Button (type: ‘button’, label: string, onClick: string), Image (type: ‘image’, src:
 285 url, width: integer), Text (type: ‘text’, content: string). Use ‘oneOf’ to ensure that a Button
 286 cannot have a ‘src’ and an Image cannot have an ‘onClick’.

287 **Required Constraints:** oneOf, const, required

288 B. SchemaGNN Hyperparameters

289 *Table 12.* Training Configuration

290 Parameter	291 Value
292 Hidden dimension	293 256
294 Attention heads	295 4
296 HGT layers	297 3
298 Dropout	299 0.1
300 Learning rate	301 1×10^{-4}
302 Batch size	303 32
304 Epochs	305 50
306 Global loss weight (λ_g)	307 1.0
308 Local loss weight (λ_l)	309 0.5

310 C. Feedback Translation Templates

311 The translator uses node-type-specific templates:

- 312 • **REF:** “This reference may point to a non-existent or incorrectly named definition. Verify that the \$ref target exists.”
- 313 • **ARRAY:** “This array definition may have invalid items schema or constraint conflicts. Verify minItems/maxItems are
 314 consistent.”
- 315 • **LOGIC:** “This logical operator may have inconsistent options. Ensure all options are valid and don’t conflict.”
- 316 • **OBJECT:** “This object may have structural issues. Check that required properties are defined.”